

# RF + Boosting Songs

February 11, 2024

## 0.1 Predicting Song Popularities with RF & XGBoost

For this assignment, we are tasked with building a classification model that can predict a track's popularity into binned categorical variables based on many features.

Here are the steps we followed in this analysis:

- Data Exploration
- Data Processing
- Random Forest model
- XGBoost model

```
[1]: # import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import plotly.graph_objects as go
from plotly.subplots import make_subplots
import seaborn as sns
from scipy import stats
import math
import ast

from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.ensemble import RandomForestClassifier
from sklearn.cluster import KMeans
from sklearn.metrics import roc_auc_score, accuracy_score, f1_score, \
    confusion_matrix, classification_report, roc_curve, auc

import xgboost as xgb
from xgboost import XGBClassifier

import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
```

## 0.2 Data Exploration

```
[2]: songs = pd.read_excel("Songs_2024.xlsx", sheet_name='spotify songs')
songs.head()
```

```
[2]:   year      track_name  track_popularity  \
0  2000           Yellow                91
1  2000  All The Small Things            84
2  2000           Breathe             69
3  2000       In the End             88
4  2000       Bye Bye Bye             74

      album  artist_name  \
0    Parachutes    Coldplay
1  Enema Of The State  blink-182
2    Breathe    Faith Hill
3  Hybrid Theory (Bonus Edition)  Linkin Park
4    No Strings Attached      *NSYNC

      artist_genres  artist_popularity  \
0  ['permanent wave'; 'pop']          86
1  ['alternative metal'; 'modern rock'; 'pop punk...  75
2  ['contemporary country'; 'country'; 'country d...  61
3  ['alternative metal'; 'nu metal'; 'post-grunge...  83
4  ['boy band'; 'dance pop'; 'pop']      65

      danceability  energy  key  loudness  mode  speechiness  acousticness  \
0          0.429  0.661  11.0      85.0   1.0         0.0281         0.00239
1          0.434  0.897   0.0      69.0   1.0         0.0488         0.01030
2          0.529  0.496   7.0      95.0   1.0         0.0290         0.17300
3          0.556  0.864   3.0      76.0   0.0         0.0584         0.00958
4          0.610  0.926   8.0      68.0   0.0         0.0479         0.03100

      instrumentalness  liveness  valence  tempo  duration_ms
0          0.000121    0.2340    0.285  173.372    266773.0
1          0.000000    0.6120    0.684  148.726    167067.0
2          0.000000    0.2510    0.278  136.859    250547.0
3          0.000000    0.2090    0.400  105.143    216880.0
4          0.001200    0.0821    0.861  172.638    200400.0
```

```
[3]: songs.describe()
```

```
[3]:   count      year  track_popularity  artist_popularity  danceability  \
count  2300.000000    2300.000000    2300.000000    2299.000000
mean    2011.000000     70.943478     72.869565     0.660116
std       6.634692    12.291526    12.179263     0.141137
min     2000.000000     0.000000    29.000000     0.162000
```

25%	2005.000000	66.000000	65.000000	0.572000
50%	2011.000000	72.000000	74.000000	0.671000
75%	2017.000000	79.000000	82.000000	0.759500
max	2022.000000	100.000000	100.000000	0.975000

	energy	key	loudness	mode	speechiness \
count	2299.000000	2299.000000	2299.000000	2299.000000	2299.000000
mean	0.693047	5.277947	73.006960	0.598521	0.097795
std	0.164838	3.628494	15.643937	0.490304	0.092445
min	0.051900	0.000000	-56.000000	0.000000	0.022500
25%	0.586000	2.000000	63.000000	0.000000	0.038000
50%	0.712000	5.000000	74.000000	1.000000	0.056800
75%	0.820000	8.000000	83.000000	1.000000	0.115500
max	0.999000	11.000000	132.000000	1.000000	0.576000

	acousticness	instrumentalness	liveness	valence	tempo \
count	2299.000000	2299.000000	2299.000000	2299.000000	2299.000000
mean	0.157689	0.013766	0.172618	0.535110	120.512450
std	0.203844	0.083990	0.131620	0.227821	27.617729
min	0.000013	0.000000	0.021000	0.037700	60.019000
25%	0.016500	0.000000	0.089950	0.360500	98.569500
50%	0.068900	0.000000	0.119000	0.540000	120.000000
75%	0.223000	0.000054	0.220000	0.722000	137.028000
max	0.978000	0.985000	0.843000	0.974000	210.857000

	duration_ms
count	2299.000000
mean	226033.494128
std	42063.678588
min	97393.000000
25%	200179.500000
50%	221653.000000
75%	245950.000000
max	688453.000000

```
[4]: songs.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2300 entries, 0 to 2299
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   year                  2300 non-null  int64
1   track_name            2300 non-null  object
2   track_popularity      2300 non-null  int64
3   album                 2300 non-null  object
4   artist_name           2300 non-null  object
```

```

5  artist_genres      2300 non-null  object
6  artist_popularity  2300 non-null  int64
7  danceability       2299 non-null  float64
8  energy             2299 non-null  float64
9  key                2299 non-null  float64
10 loudness           2299 non-null  float64
11 mode               2299 non-null  float64
12 speechiness        2299 non-null  float64
13 acousticness       2299 non-null  float64
14 instrumentalness    2299 non-null  float64
15 liveness           2299 non-null  float64
16 valence             2299 non-null  float64
17 tempo              2299 non-null  float64
18 duration_ms        2299 non-null  float64
dtypes: float64(12), int64(3), object(4)
memory usage: 341.5+ KB

```

```

[5]: #Checking number of unique rows in each feature
songs.nunique().sort_values()

```

```

[5]: mode                2
key                    12
year                  23
artist_popularity     62
track_popularity      71
loudness              104
danceability          585
energy                641
artist_genres         698
liveness              807
valence               827
instrumentalness      841
speechiness           860
artist_name           891
acousticness          1322
album                 1663
duration_ms           2033
tempo                 2074
track_name            2121
dtype: int64

```

```

[6]: # visualizing missing rows as numpy array for ease
missing_rows = songs[songs.isnull().any(axis=1)].to_numpy()

# Display the rows with missing values
print(missing_rows)

```

```

[[2004 'These Words' 68 'Unwritten' 'Natasha Bedingfield'

```

```
"['dance pop'; 'pop'; 'post-teen pop']" 64 nan nan nan nan nan nan nan
nan nan nan nan nan]
```

```
[7]: songs.isnull().sum()
```

```
[7]: year          0
     track_name     0
     track_popularity 0
     album          0
     artist_name    0
     artist_genres  0
     artist_popularity 0
     danceability   1
     energy         1
     key            1
     loudness       1
     mode           1
     speechiness    1
     acousticness   1
     instrumentalness 1
     liveness       1
     valence        1
     tempo          1
     duration_ms    1
     dtype: int64
```

```
[8]: # Dealing with missing values and duplicates
songs = songs.drop_duplicates(subset=['track_name', 'artist_name']) #Bare in
    ↪ mind that we can have songs with the same name but different artists, so we
    ↪ are not removing those.

for column in songs.columns:
    if songs[column].isna().any() and songs[column].dtype != 'object':
        # Fill null values with the mean
        songs[column].fillna(songs[column].mean(), inplace=True)
```

```
[9]: songs.isnull().sum()
```

```
[9]: year          0
     track_name     0
     track_popularity 0
     album          0
     artist_name    0
     artist_genres  0
     artist_popularity 0
     danceability   0
```

```

energy          0
key             0
loudness        0
mode            0
speechiness     0
acousticness    0
instrumentalness 0
liveness        0
valence         0
tempo           0
duration_ms     0
dtype: int64

```

```

[10]: # choosing the metric variables to check outlier counts
columns_to_check = [
    'year', 'track_popularity', 'artist_popularity', 'danceability', 'energy', 'key', 'loudness', 'mo

def detect_outliers_iqr(data, column):
    """
    Detects outliers in a dataframe column based on the IQR method.

    Parameters:
    - data: pandas DataFrame.
    - column: String name of the column to check for outliers.

    Returns:
    - DataFrame containing the outliers.
    """
    Q1 = data[column].quantile(0.25)
    Q3 = data[column].quantile(0.75)
    IQR = Q3 - Q1
    # 5 is a high value to check for drastic outliers
    lower_bound = Q1 - 5 * IQR
    upper_bound = Q3 + 5 * IQR
    return data[(data[column] < lower_bound) | (data[column] > upper_bound)]

def detect_outliers_zscore(data, column, threshold=3):
    """
    Detects outliers in a dataframe column based on the Z-score method.

    Parameters:
    - data: pandas DataFrame.
    - column: String name of the column to check for outliers.
    - threshold: Z-score threshold for considering a point as an outlier.

    Returns:
    - DataFrame containing the outliers.

```

```

"""
z_scores = stats.zscore(data[column])
return data[abs(z_scores) > threshold]

outliers_count_z = {}
outliers_count_iqr = {}

# Looping through the columns to check for outliers
for column in columns_to_check:
    outliers_iqr = detect_outliers_iqr(songs, column)
    outliers_z = detect_outliers_zscore(songs, column)
    outliers_count_z[column] = len(outliers_iqr)
    outliers_count_iqr[column] = len(outliers_z)

# Creating DataFrames for Z-score and IQR outlier counts
df_outliers_z = pd.DataFrame(list(outliers_count_z.items()), columns=['Column',
↪ 'Outliers_Z'])
df_outliers_iqr = pd.DataFrame(list(outliers_count_iqr.items()),
↪ columns=['Column', 'Outliers_IQR'])

# Merging the two DataFrames on the 'Column' column
df_outliers = pd.merge(df_outliers_z, df_outliers_iqr, on='Column')

print(df_outliers)

```

	Column	Outliers_Z	Outliers_IQR
0	year	0	0
1	track_popularity	0	38
2	artist_popularity	0	7
3	danceability	0	9
4	energy	0	13
5	key	0	0
6	loudness	1	16
7	mode	0	0
8	speechiness	2	40
9	acousticness	0	56
10	instrumentalness	375	36
11	liveness	0	60
12	valence	0	0
13	tempo	0	6
14	duration_ms	3	17

```

[11]: # Replace semicolons with commas in the 'artist_genres' column before applying
↪ ast.literal_eval
songs['artist_genres'] = songs['artist_genres'].apply(lambda x: x.replace(';',
↪ ',') if isinstance(x, str) else x)

```

```

# Now safely apply ast.literal_eval to convert string representations of lists
# into actual lists
songs['artist_genres'] = songs['artist_genres'].apply(lambda x: ast.
    literal_eval(x) if isinstance(x, str) else x)

songs = songs.dropna(subset=['artist_genres'])

# Flatten the list of genres from all rows into a single list
all_genres = sum(songs['artist_genres'], [])

# Convert the list to a set to remove duplicates, then count the unique elements
unique_genres = set(all_genres)
number_of_unique_genres = len(unique_genres)

print(f"There are {number_of_unique_genres} unique genres in the dataset.")

```

There are 437 unique genres in the dataset.

### 0.2.1 Experimenting with Genre

```

[12]: # Normalize genre names: lowercase and remove special characters if necessary
songs['artist_genres'] = ["".join(x).lower().replace('[^a-z\s]', '') for x in
    songs['artist_genres']]

# Initialize TF-IDF Vectorizer
tfidf = TfidfVectorizer(max_features=1000)

# Fit and transform the genre names
genre_tfidf = tfidf.fit_transform(songs['artist_genres'])

# Apply K-Means clustering
kmeans = KMeans(n_clusters=10, random_state=42)
songs['genre_cluster'] = kmeans.fit_predict(genre_tfidf)

# Display the number of genres per cluster to check the distribution
print(songs['genre_cluster'].value_counts())

```

```

genre_cluster
2    946
7    260
1    238
8    168
3    154
9    132
0    107
5     89
6     49

```



4 38

Name: count, dtype: int64

```
[13]: # For simplicity, let's see some examples of genres in each cluster
for cluster in sorted(songs['genre_cluster'].unique()):
    print(f"Cluster {cluster}:")
    examples = songs[songs['genre_cluster'] == cluster]['artist_genres'].
    ↪unique()
    print(examples[:5]) # Print 5 examples per cluster
```

Cluster 0:

```
['detroit hip hophip hoprap'
 'east coast hip hopgangster raphip hoppop raprap'
 'east coast hip hopgangster raphardcore hip hophip hoprap'
 'detroit hip hopgangster raphip hop'
 'dance popeast coast hip hopgangster raphip hophip poppop rapqueens hip
 hoprapurban contemporary']
```

Cluster 1:

```
['pop rock' 'canadian pop' 'dance pop' 'hip pop' 'pop']
```

Cluster 2:

```
['permanent wavepop' 'boy banddance poppop' 'dance rockeuropop'
 'disco housefilter house' 'bouncy houseeurodance']
```

Cluster 3:

```
['dance poppop' 'canadian poppop' 'hip poppop' 'poppop r&b'
 'dance poppop danceuk dance']
```

Cluster 4:

```
['popsinger-songwriter popuk pop' 'popuk pop' 'popuk danceuk funky']
```

Cluster 5:

```
['g funkgangster raphip hoprapwest coast rap' 'comedy rap' 'rap rock'
 'dance popdancehallpoppop rap' 'dance poppoppop rap']
```

Cluster 6:

```
['contemporary r&bhip popr&b'
 'dance popr&bsouthern hip hopurban contemporary' 'popr&b'
 'hip popr&bsouthern hip hopurban contemporary'
 'contemporary r&bdirty south raphip popr&bsouthern hip hopurban contemporary']
```

Cluster 7:

```
['east coast hip hophardcore hip hop'
 'atl hip hopdirty south raphip hopold school atlanta hip hoprapsouthern hip
 hop'
 'atl hip hopcontemporary r&bdance poppopr&brapsouth carolina hip hopurban
 contemporary'
 'dance pophip hophip popneo soulpop rapr&brapurban contemporaryvirginia hip
 hop'
 'hip popurban contemporary']
```

Cluster 8:

```
['contemporary countrycountrycountry dawncountry road'
 'contemporary r&bdirty south raphip popr&burban contemporary'
 'dance popgirl grouppopr&burban contemporary']
```

```

'contemporary countrycountrycountry road'
'dance popgangster raphip hoppop raprapst louis rapurban contemporary']
Cluster 9:

```

```

['alternative metalmodern rockpop punkpunkrocksocial pop punk'
'alternative metalnu metalpost-grungerap metalrock'
'alternative rockfunk metalfunk rockpermanent waverock'
'alternative metalnu metalpost-grungerock'
'alternative metalfunk metalnu metalpost-grungerap metalrock']

```

```

[14]: cluster_labels = {
      0: 'Urban Hip Hop',
      1: 'Mainstream Pop',
      2: 'Dance/Electronic',
      3: 'Contemporary R&B',
      4: 'UK Pop/Funk',
      5: 'Diverse Hip Hop',
      6: 'Modern R&B',
      7: 'Hip Hop/Rap Fusion',
      8: 'Country',
      9: 'Alt Rock/Metal Fusion'
    }

```

```

[15]: # Map the cluster labels to the dataset
songs['genre_label'] = songs['genre_cluster'].map(cluster_labels)
print(songs[['artist_genres', 'genre_cluster', 'genre_label']].head(10))

```

	artist_genres	genre_cluster	\
0	permanent wavepop	2	
1	alternative metalmodern rockpop punkpunkrockso...	9	
2	contemporary countrycountrycountry dawncountry...	8	
3	alternative metalnu metalpost-grungerap metalrock	9	
4	boy banddance poppop	2	
5	contemporary r&bdirty south raphip popr&burban...	8	
6	detroit hip hophip hoprap	0	
7	dance rockeuropop	2	
8	dance popgirl grouppopr&burban contemporary	8	
9	alternative rockfunk metalfunk rockpermanent w...	9	

	genre_label
0	Dance/Electronic
1	Alt Rock/Metal Fusion
2	Country
3	Alt Rock/Metal Fusion
4	Dance/Electronic
5	Country
6	Urban Hip Hop
7	Dance/Electronic
8	Country

## 9 Alt Rock/Metal Fusion

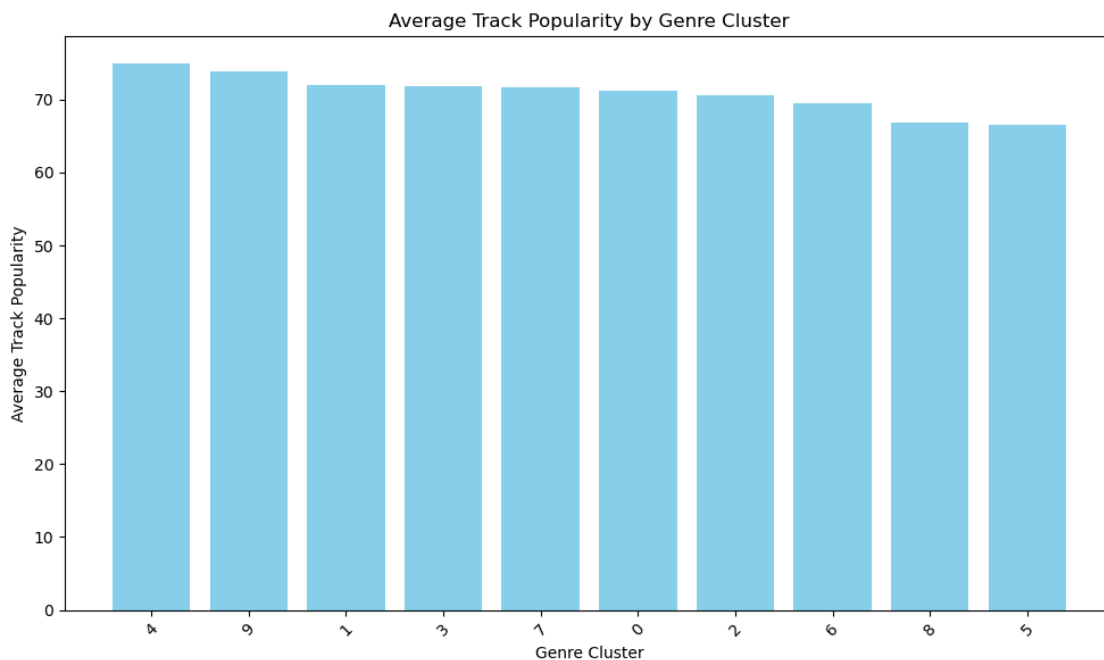
```
[16]: songs.to_excel('genre_clustered.xlsx', index=False)
```

```
[17]: # Load the data
songs = pd.read_excel('genre_clustered.xlsx' )

# Calculate average track popularity for each genre cluster
avg_popularity_by_cluster = songs.groupby('genre_cluster')['track_popularity'].
    ↪mean().reset_index()

# Sort the clusters for better visualization
avg_popularity_by_cluster = avg_popularity_by_cluster.
    ↪sort_values(by='track_popularity', ascending=False)

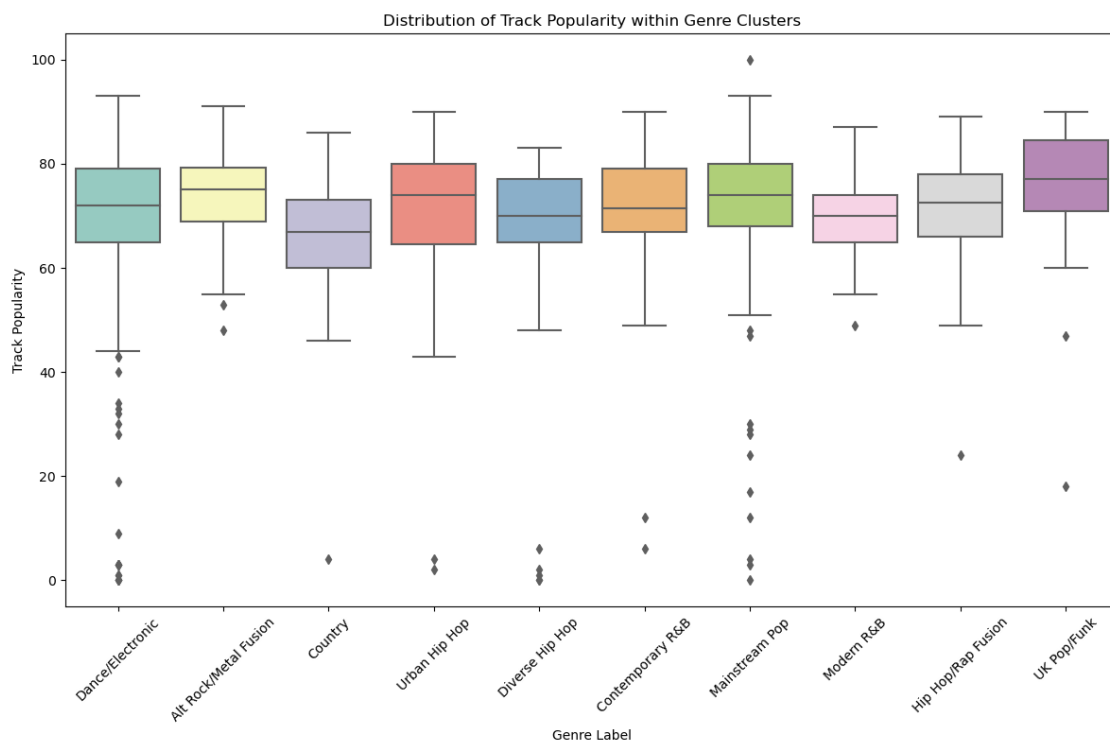
# Plotting
plt.figure(figsize=(10, 6))
plt.bar(avg_popularity_by_cluster['genre_cluster'].astype(str),
    ↪avg_popularity_by_cluster['track_popularity'], color='skyblue')
plt.xlabel('Genre Cluster')
plt.ylabel('Average Track Popularity')
plt.title('Average Track Popularity by Genre Cluster')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
[18]: # Creating the boxplot
plt.figure(figsize=(12, 8))
sns.boxplot(x='genre_label', y='track_popularity', data=songs, palette="Set3")

plt.title('Distribution of Track Popularity within Genre Clusters')
plt.xlabel('Genre Label')
plt.ylabel('Track Popularity')
plt.xticks(rotation=45) # Rotates the labels on the x-axis for better
    ↪ readability

plt.tight_layout()
plt.show()
```



```
[19]: # list of your numerical feature column names
nf = ['danceability', 'energy', 'loudness', 'speechiness', 'acousticness',
    ↪ 'instrumentalness', 'liveness', 'valence', 'tempo', 'duration_ms']

n = 5 # Columns per row
rows = math.ceil(len(nf)/n) # Number of rows needed

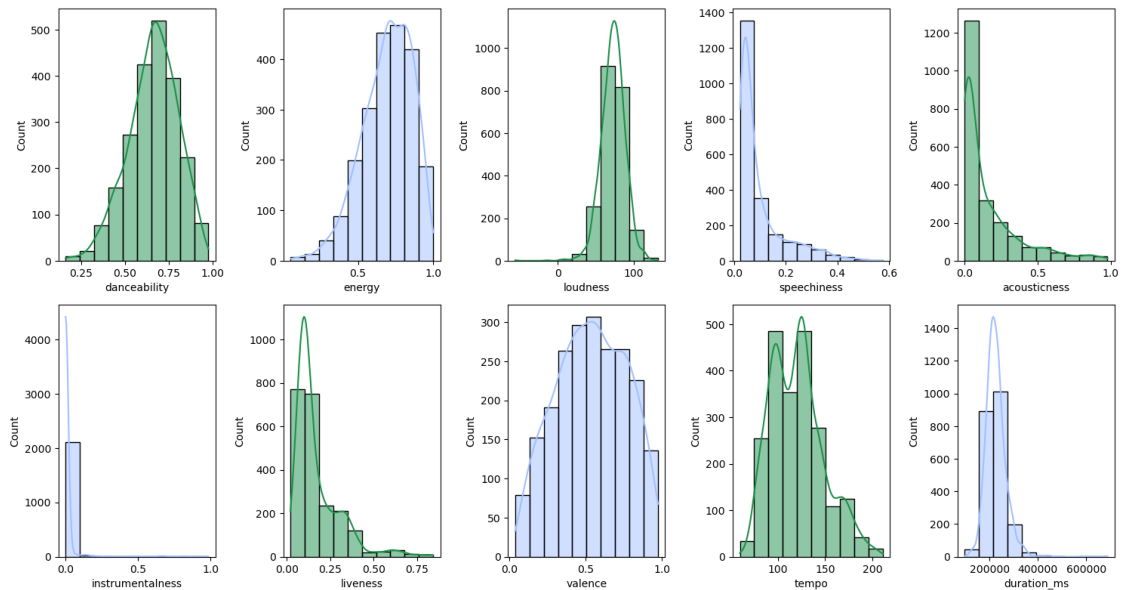
# Define colors to cycle between green and blue
clr = ['#1E914C', '#A2C1FF'] # using specific hex codes
```

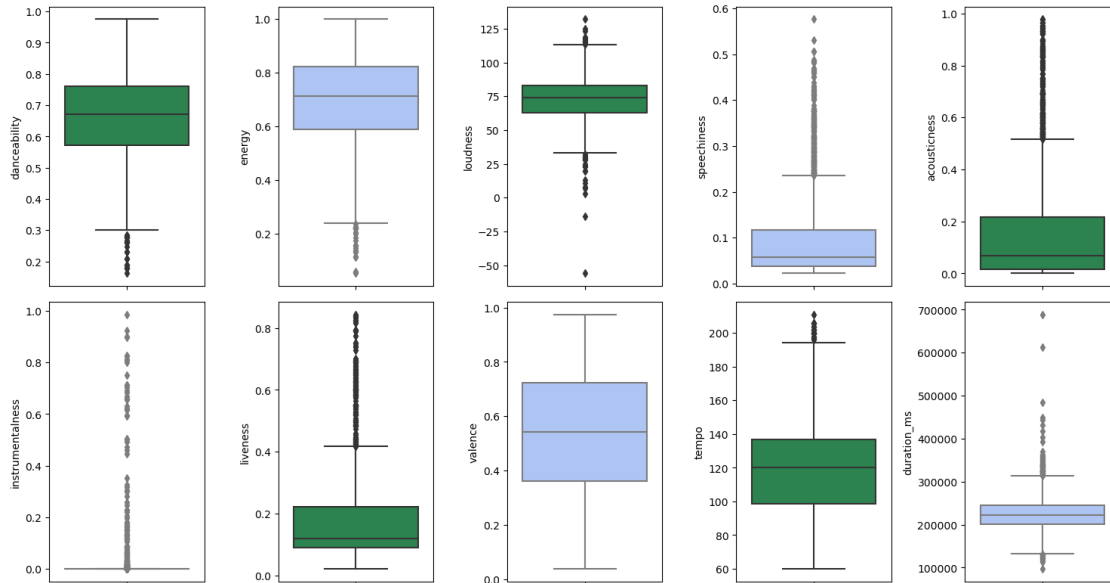
```

# Distribution plots
plt.figure(figsize=[15, 4*rows])
for i, feature in enumerate(nf):
    plt.subplot(rows, n, i+1)
    sns.histplot(songs[feature], kde=True, color=clr[i % len(clr)],
        edgecolor="black", bins=10)
plt.tight_layout()
plt.show()

# Box plots
plt.figure(figsize=[15, 4*rows])
for i, feature in enumerate(nf):
    plt.subplot(rows, n, i+1)
    sns.boxplot(y=songs[feature], color=clr[i % len(clr)])
plt.tight_layout()
plt.show()

```





```
[20]: # Create a subplot grid: 4 rows, 2 columns, making it more compact
fig = make_subplots(rows=4, cols=2, subplot_titles=(
    'Artist Popularity over Years', 'Artist Popularity vs Track Popularity',
    'Valence vs Danceability', 'Energy vs Valence',
    'Danceability vs Energy', 'Loudness vs Acousticness (Negative Correlation)',
    'Energy vs Acousticness (Negative Correlation)'
))

# Add each scatter plot to the grid, adjust the row and col accordingly
fig.add_trace(go.Scatter(x=songs['year'], y=songs['artist_popularity'],
    ↪mode='markers', marker=dict(color='#1E914C')), row=1, col=1)
fig.add_trace(go.Scatter(x=songs['artist_popularity'],
    ↪y=songs['track_popularity'], mode='markers', marker=dict(color='#A2C1FF')),
    ↪row=1, col=2)
fig.add_trace(go.Scatter(x=songs['valence'], y=songs['danceability'],
    ↪mode='markers', marker=dict(color='#1E914C')), row=2, col=1)
fig.add_trace(go.Scatter(x=songs['energy'], y=songs['valence'], mode='markers',
    ↪marker=dict(color='#A2C1FF')), row=2, col=2)
fig.add_trace(go.Scatter(x=songs['danceability'], y=songs['energy'],
    ↪mode='markers', marker=dict(color='#1E914C')), row=3, col=1)
fig.add_trace(go.Scatter(x=songs['loudness'], y=songs['acousticness'],
    ↪mode='markers', marker=dict(color='#A2C1FF')), row=3, col=2)
fig.add_trace(go.Scatter(x=songs['energy'], y=songs['acousticness'],
    ↪mode='markers', marker=dict(color='#1E914C')), row=4, col=1)

# Since we have an odd number of plots, one subplot will be empty
```

```
# Update layout for a cleaner look
fig.update_layout(height=1200, width=800, title_text="Scatter Plots Grid",
↳showlegend=False)
fig.show()
```

```
[21]: # Set the figure size
plt.figure(figsize=[8, 4])
# Creating graphic with seaborn using the 'track_popularity' column
sns.distplot(songs['track_popularity'], color='g', kde_kws={"color": "k", "lw": 2}, hist_kws={"edgecolor": "black", "linewidth": 1}, bins=50)
# Establecer el título del gráfico
```

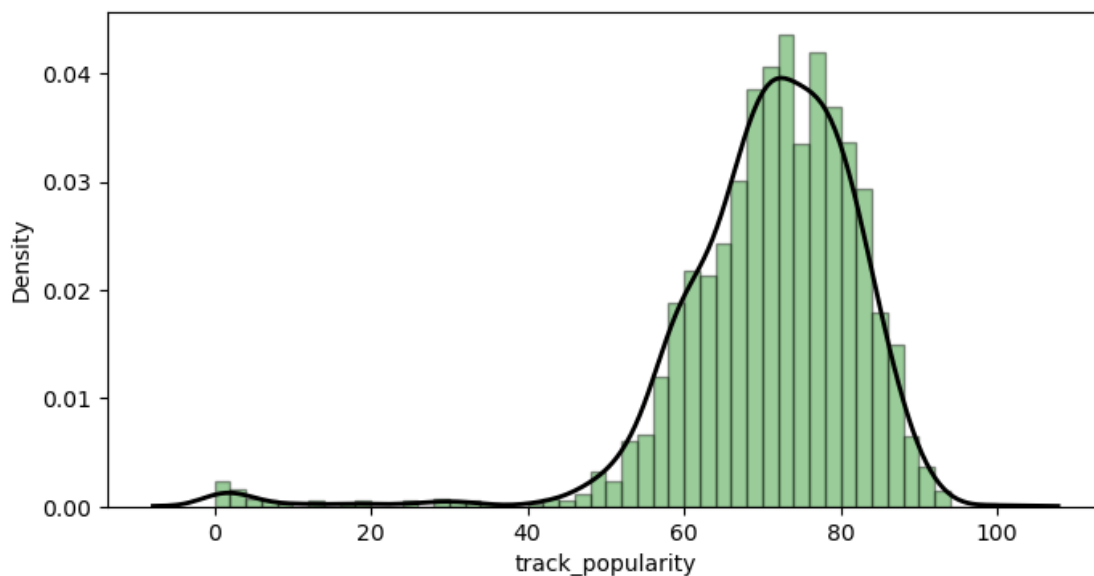
/var/folders/r8/9knv8vvj3m98jv1qltst72xr0000gn/T/ipykernel\_72994/1722633087.py:4  
: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

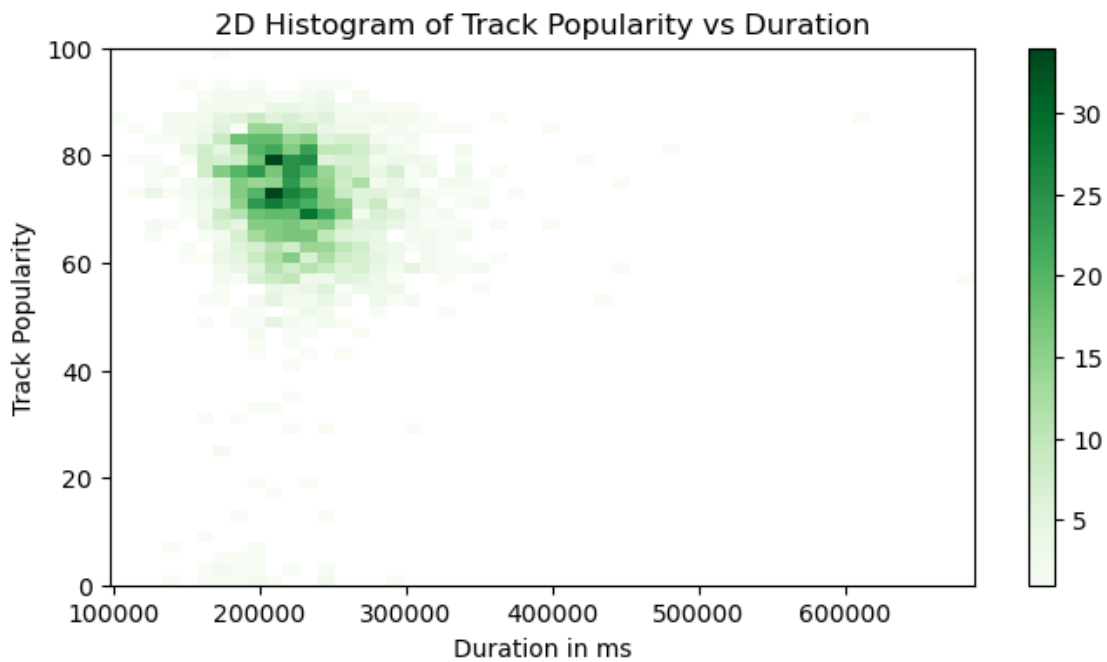
Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
[21]: <Axes: xlabel='track_popularity', ylabel='Density'>
```

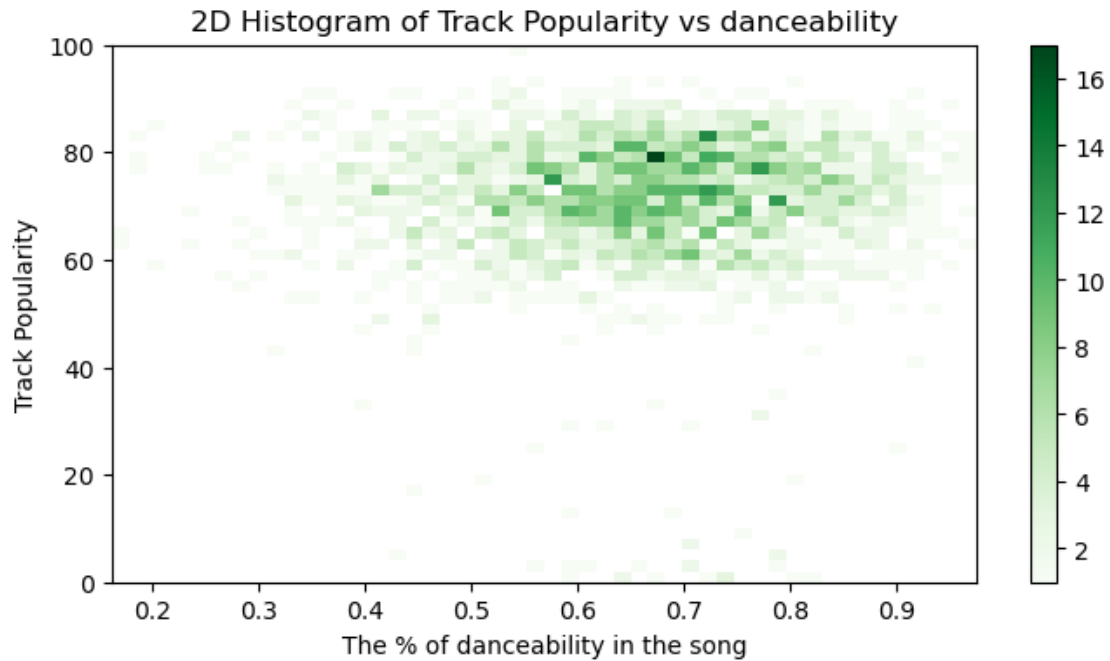


```
[22]: plt.figure(figsize=[8, 4])
plt.hist2d(songs['duration_ms'], songs['track_popularity'], bins=[50, 50],
           cmap='Greens', cmin=1)
plt.xlabel('Duration in ms')
plt.ylabel('Track Popularity')
plt.title('2D Histogram of Track Popularity vs Duration')
plt.colorbar()
plt.show()
```

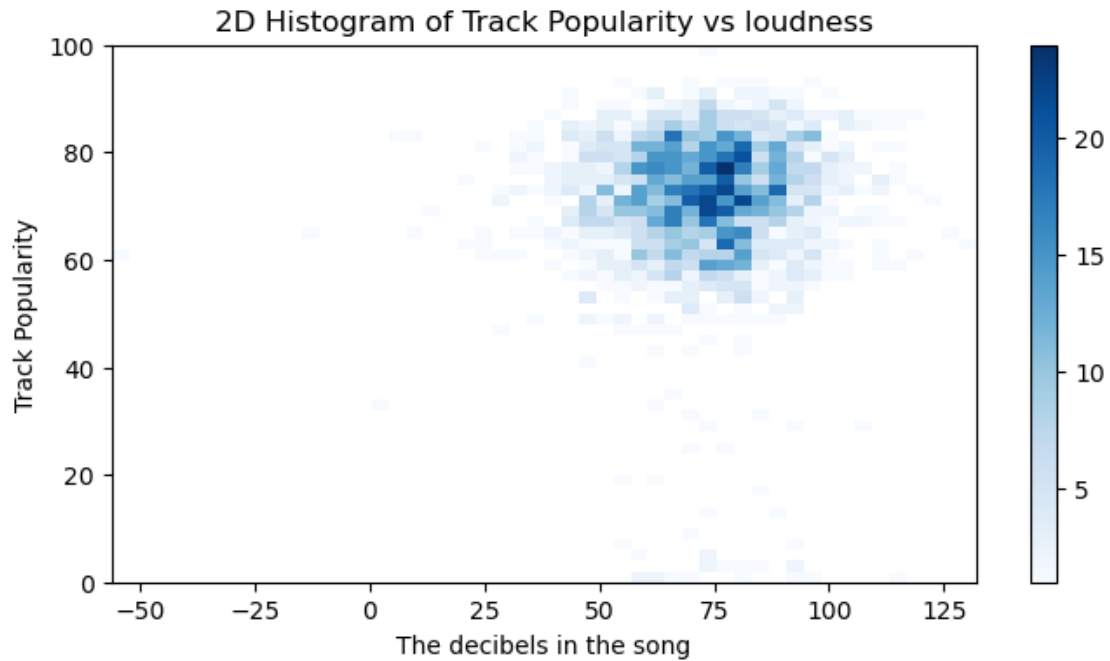


```
[23]: plt.figure(figsize=[8, 4])
plt.hist2d(songs['danceability'], songs['track_popularity'], bins=[50, 50],
           cmap='Greens', cmin=1)
plt.xlabel('The % of danceability in the song')
plt.ylabel('Track Popularity')
plt.title('2D Histogram of Track Popularity vs danceability')
plt.colorbar()
plt.show()
```



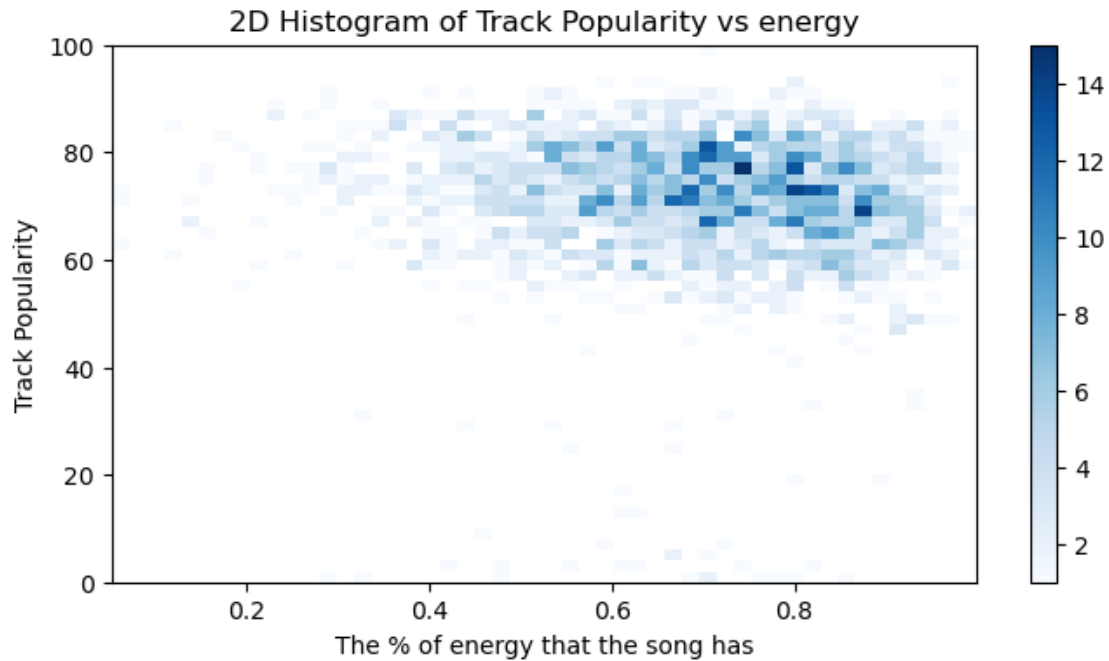


```
[24]: # Ajustar el tamaño de la figura con matplotlib
plt.figure(figsize=[8, 4])
plt.hist2d(songs['loudness'], songs['track_popularity'], bins=[50, 50],
           cmap='Blues', cmin=1)
plt.xlabel('The decibels in the song')
plt.ylabel('Track Popularity')
plt.title('2D Histogram of Track Popularity vs loudness')
plt.colorbar()
plt.show()
```



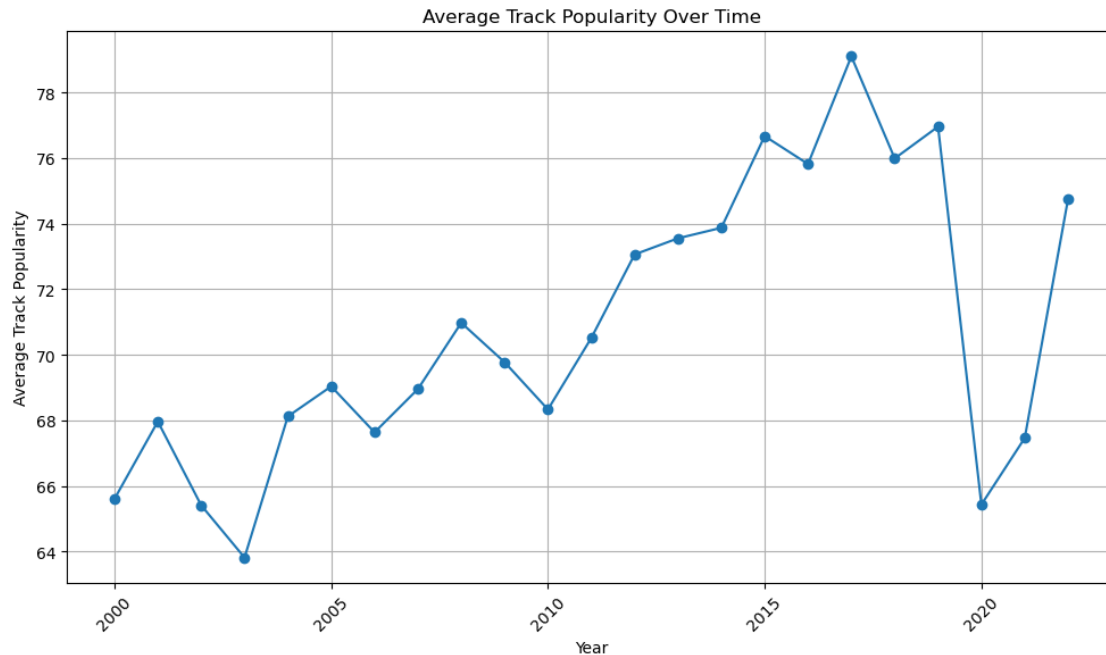
```
[25]: # Ajustar el tamaño de la figura con matplotlib
plt.figure(figsize=[8, 4])

plt.hist2d(songs['energy'], songs['track_popularity'], bins=[50, 50],
           cmap='Blues', cmin=1)
plt.xlabel('The % of energy that the song has')
plt.ylabel('Track Popularity')
plt.title('2D Histogram of Track Popularity vs energy')
plt.colorbar()
plt.show()
```



```
[26]: # Step 2: Group the data by year and calculate the average track popularity
average_popularity_by_year = songs.groupby('year')['track_popularity'].mean()

# Step 3: Plot the time series
plt.figure(figsize=(10, 6))
average_popularity_by_year.plot(kind='line', marker='o', linestyle='--')
plt.title('Average Track Popularity Over Time')
plt.xlabel('Year')
plt.ylabel('Average Track Popularity')
plt.grid(True)
plt.xticks(rotation=45) # Rotate x-axis labels for better readability
plt.tight_layout() # Adjust layout to make room for the rotated x-axis labels
plt.show()
```



## 0.3 Data Processing

We are going to follow the following steps for our data processing: 1. Feature Engineering 2. Binning target variable

### 0.3.1 Feature Engineering

```
[27]: # New feature called Melodic Energy that captures the interaction between
      ↪energy and danceability
songs['Melodic Energy'] = songs['energy']*songs['danceability']

# New feature called Emotional Tempo that captures the interaction between
      ↪valence and tempo
songs['Emotional Tempo'] = songs['valence']*songs['tempo']
# weight = 0.5
# songs['Acoustic Instruments'] = (weight * songs['instrumentalness']) +
      ↪(weight * songs['acousticness']) + 1

# songs['Loudness per ms'] = songs['loudness']/songs['duration_ms']

# Dropping the features that were used to create the new features to avoid
      ↪unnecessary multicollinearity
songs.drop(['energy', 'danceability', 'valence', 'tempo'], axis=1, inplace=True)
```

```
[28]: # Select a subset of popular artists based on a threshold of artist popularity
      ↪for the median
high_popularity_artists = songs[songs['artist_popularity'] >
      ↪songs['artist_popularity'].median()]
# Calculate the median number of songs per popular artist
song_count_popular_artists = high_popularity_artists.groupby('artist_name').
      ↪size().median()
# Calculate the total number of songs
total_songs_count = len(songs)
# Calculate the ratio of songs by popular artists to total songs
high_popularity_ratio = song_count_popular_artists / total_songs_count

# Calculating individual songs per artist
songs_per_artist = songs.groupby('artist_name').size()
# Calculate artist ratio, handling potential division by zero errors
artist_ratio = songs_per_artist.div(total_songs_count, fill_value=0)
# Map the artist ratio to the original dataset
songs['artist_ratio'] = songs['artist_name'].map(artist_ratio)

# Calculating the difference between the artist ratio and the high popularity
      ↪artist ratio
songs['popularity_difference'] = songs['artist_ratio'] - high_popularity_ratio
# Dropping the artist ratio column to avoid multicollinearity
songs.drop('artist_ratio', axis=1, inplace=True)

[29]: ## Creating encoded binary dummy columns for genre_label column
# genre_dummies = pd.get_dummies(songs['genre_label'], drop_first=True)
## Concatenating the dummy columns to the original dataset
# songs = pd.concat([songs, genre_dummies], axis=1)

[30]: # Drop the non-numerical columns where we don't find meaning
songs = songs.drop(columns=['track_name', 'album', 'artist_name',
      ↪'artist_genres', 'key', 'mode', 'genre_cluster', 'genre_label'], axis=1)
```

### 0.3.2 Binning

We decided to go with a binning approach into 2 bins, splitting on the track popularity of 70. From exploring the data, we saw that 70 provides us a pretty equal split.

```
[31]: songs['popularity_binned'] = pd.cut(songs['track_popularity'], bins=[0,70,100],
      ↪labels=['Work in Progress', 'Hit'])
songs['popularity_binned'] = songs['popularity_binned'].
      ↪fillna(songs['popularity_binned'].mode()[0])

[32]: # counting the number of records in our sign category
print(songs['popularity_binned'].value_counts())
```

```

popularity_binned
Hit          1250
Work in Progress  931
Name: count, dtype: int64

```

```
[33]: # Export cleaned version of the dataset to an Excel file to save the changes
songs.to_excel('Songs_2024_cleaned.xlsx', index=False)
```

```
[34]: data=pd.read_excel("Songs_2024_cleaned.xlsx")
data.head()
```

```
[34]:
```

	year	track_popularity	artist_popularity	loudness	speechiness	\
0	2000	91	86	85.0	0.0281	
1	2000	84	75	69.0	0.0488	
2	2000	69	61	95.0	0.0290	
3	2000	88	83	76.0	0.0584	
4	2000	74	65	68.0	0.0479	

	acousticness	instrumentalness	liveness	duration_ms	Melodic Energy	\
0	0.00239	0.000121	0.2340	266773.0	0.283569	
1	0.01030	0.000000	0.6120	167067.0	0.389298	
2	0.17300	0.000000	0.2510	250547.0	0.262384	
3	0.00958	0.000000	0.2090	216880.0	0.480384	
4	0.03100	0.001200	0.0821	200400.0	0.564860	

	Emotional Tempo	popularity_difference	popularity_binned
0	49.411020	0.005961	Hit
1	101.728584	-0.000459	Hit
2	38.046802	-0.000459	Work in Progress
3	42.057200	0.002751	Hit
4	148.641318	0.000917	Hit

```
[35]: data.describe()
```

```
[35]:
```

	year	track_popularity	artist_popularity	loudness	\
count	2181.000000	2181.000000	2181.000000	2181.000000	
mean	2010.862907	70.775332	72.696470	72.957339	
std	6.632351	12.114627	12.204621	15.709096	
min	2000.000000	0.000000	29.000000	-56.000000	
25%	2005.000000	65.000000	65.000000	63.000000	
50%	2011.000000	72.000000	74.000000	74.000000	
75%	2017.000000	78.000000	82.000000	83.000000	
max	2022.000000	100.000000	100.000000	132.000000	

	speechiness	acousticness	instrumentalness	liveness	\
count	2181.000000	2181.000000	2181.000000	2181.000000	
mean	0.098194	0.155747	0.014084	0.173459	

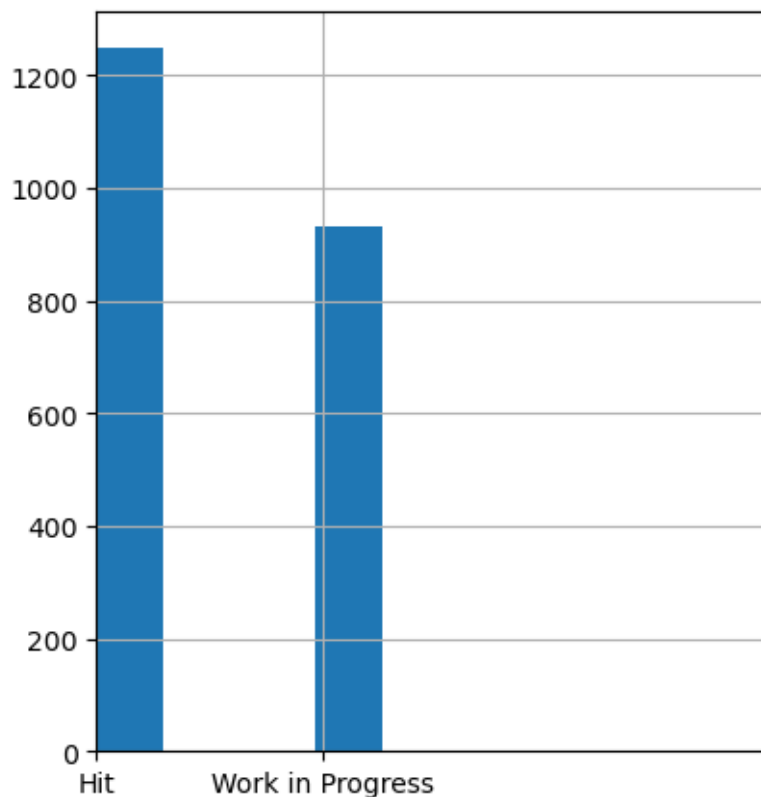
std	0.092620	0.202922	0.085464	0.132511
min	0.022500	0.000013	0.000000	0.021000
25%	0.038200	0.016400	0.000000	0.090000
50%	0.057100	0.068500	0.000000	0.120000
75%	0.117000	0.216000	0.000054	0.221000
max	0.576000	0.978000	0.985000	0.843000

	duration_ms	Melodic Energy	Emotional Tempo	popularity_difference
count	2181.000000	2181.000000	2181.000000	2181.000000
mean	226263.307339	0.457684	64.575285	0.001673
std	42192.772559	0.138661	31.459754	0.003133
min	97393.000000	0.020124	3.656558	-0.000917
25%	200360.000000	0.366704	40.797280	-0.000459
50%	221840.000000	0.465010	62.314008	0.000459
75%	245940.000000	0.556206	85.227034	0.002751
max	688453.000000	0.848216	172.636704	0.012380

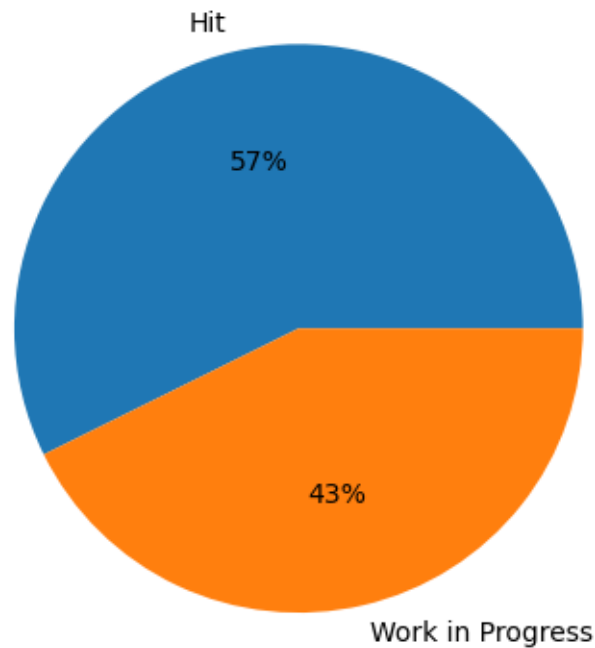
```
[36]: # Exploring the binning of the target variable
data['popularity_binned'].hist(bins=30, figsize=(10, 5), width=0.3, ax=plt.
    ↳ subplot(1, 2, 1))

plt.xlim(0, 3)
```

[36]: (0.0, 3.0)



```
[37]: unique, counts = np.unique(data["popularity_binned"], return_counts=True)
plt.pie(counts, labels=unique, autopct='%0f%%');
```



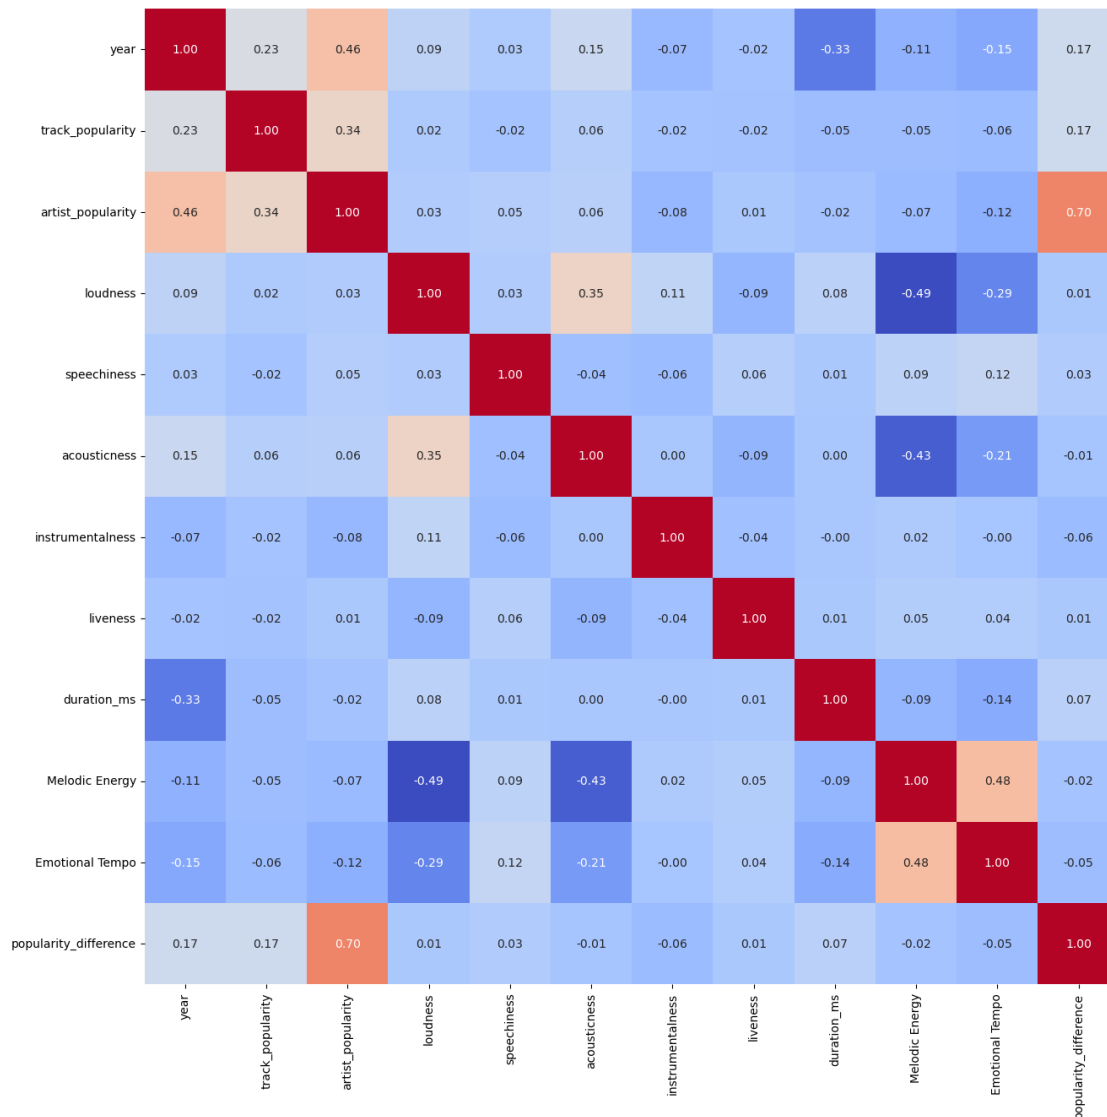
```
[38]: data.columns
```

```
[38]: Index(['year', 'track_popularity', 'artist_popularity', 'loudness',
            'speechiness', 'acousticness', 'instrumentalness', 'liveness',
            'duration_ms', 'Melodic Energy', 'Emotional Tempo',
            'popularity_difference', 'popularity_binned'],
           dtype='object')
```

```
[39]: # Exploring the correlations between all features, both old and engineered new
      ↪ ones
plt.figure(figsize=(15, 15))
to_plot_corr = data.drop(['popularity_binned'], axis=1)
sns.heatmap(to_plot_corr.corr(), annot=True, fmt=".2f", cmap='coolwarm',
      ↪ cbar=False, square=True)
```

```
[39]: <Axes: >
```





## Splitting the data for models

```
[40]: # Splitting the data into train and test

# Separate features from response field
X = data.drop(['popularity_binned', 'track_popularity'], axis=1)
y = data['popularity_binned']

# This encodes the target variable, y, into 0 and 1 (1 meaning Hit and 0
↳ meaning Work in Progress)
encode = LabelEncoder()
y = encode.fit_transform(y)
```

```
# train test split with 20% testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =0.2,
↳random_state= 123)
```

```
[41]: y_train.shape, y_test.shape
```

```
[41]: ((1744,), (437,))
```

## 0.4 Random Forest model

We try first with Random Forest. This model has some features for consideration: - Number of trees (estimators) does not impact overfitting - Max features matters here (just bagging with number of features used for model limited) - Must tune to find best hyperparameters: - 'max\_depth': maximum number of levels in each decision tree - 'min\_samples\_leaf': minimum number of samples required to be at a leaf node - 'max\_features': max features considered for each decision tree (we have 11 features total)

Potential scoring metrics: - Accuracy:  $(TP + TN) / (TP + TN + FP + FN)$  - f1 score =  $2 * (\text{precision} * \text{recall}) / (\text{precision} + \text{recall})$  - roc\_auc = Area under the ROC curve

Given our business model, we will choose the best metric to be the F1 score, given as it captures the interaction between precision and recall quite nicely.

```
[42]: # Setting the parameter for GridSearchCV to focus on accuracy score optimization
scorer = 'accuracy'
```

```
[43]: # Instantiate the RF model
rf = RandomForestClassifier(n_estimators=400, oob_score=True, random_state=123)
rf.fit(X_train, y_train)
```

```
[43]: RandomForestClassifier(n_estimators=400, oob_score=True, random_state=123)
```

Now we do a grid searching of best hyper parameters for Random Forest.

```
[44]: # grid search for best RF parameters
rf_params = {
    'max_depth': [4,5,6], # maximum number of levels in each decision tree
    'min_samples_leaf': [1,3,5,7], # minimum number of samples required to be
↳at a leaf node
    'max_features': [8,9,10,11] # we have 10 features
}
# Instantiate the grid search model with the RF model and the parameter grid
grid_rf = GridSearchCV(estimator = rf,
                        param_grid = rf_params,
                        cv=5,
                        scoring=scorer,
                        verbose=2)
# Fit the grid search to the data
```

```
grid_rf.fit(X_train, y_train)
```

Fitting 5 folds for each of 48 candidates, totalling 240 fits

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

```
[CV] END ...max_depth=6, max_features=11, min_samples_leaf=7; total time= 2.3s
[CV] END ...max_depth=6, max_features=11, min_samples_leaf=7; total time= 2.7s
[CV] END ...max_depth=6, max_features=11, min_samples_leaf=7; total time= 2.4s
[CV] END ...max_depth=6, max_features=11, min_samples_leaf=7; total time= 2.4s
```

```
[44]: GridSearchCV(cv=5,
                  estimator=RandomForestClassifier(n_estimators=400, oob_score=True,
                                                    random_state=123),
                  param_grid={'max_depth': [4, 5, 6], 'max_features': [8, 9, 10, 11],
                              'min_samples_leaf': [1, 3, 5, 7]},
                  scoring='accuracy', verbose=2)
```

```
[45]: # Random Forest model with the best parameters found by GridSearchCV
random_forest_model = RandomForestClassifier(
    n_estimators = 400, random_state=123, #
    max_depth=grid_rf.best_params_['max_depth'], # from our grid search
    max_features=grid_rf.best_params_['max_features'], # from our grid search
    min_samples_leaf=grid_rf.best_params_['min_samples_leaf'], # from our grid_
    ↪search
    verbose=2) # verbose = 2 means that we want to see the output

# Fit the hyperparameter-tuned model to the training data
random_forest_model.fit(X_train, y_train)
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.0s remaining: 0.0s
```

```
building tree 1 of 400
building tree 2 of 400
building tree 3 of 400
building tree 4 of 400
building tree 5 of 400
building tree 6 of 400
building tree 7 of 400
building tree 8 of 400
building tree 9 of 400
building tree 10 of 400
building tree 11 of 400
building tree 12 of 400
building tree 13 of 400
building tree 14 of 400
building tree 15 of 400
building tree 16 of 400
building tree 17 of 400
building tree 18 of 400
building tree 19 of 400
building tree 20 of 400
building tree 21 of 400
```



building tree 22 of 400  
building tree 23 of 400  
building tree 24 of 400  
building tree 25 of 400  
building tree 26 of 400  
building tree 27 of 400  
building tree 28 of 400  
building tree 29 of 400  
building tree 30 of 400  
building tree 31 of 400  
building tree 32 of 400  
building tree 33 of 400  
building tree 34 of 400  
building tree 35 of 400  
building tree 36 of 400  
building tree 37 of 400  
building tree 38 of 400  
building tree 39 of 400  
building tree 40 of 400  
building tree 41 of 400  
building tree 42 of 400  
building tree 43 of 400  
building tree 44 of 400  
building tree 45 of 400  
building tree 46 of 400  
building tree 47 of 400  
building tree 48 of 400  
building tree 49 of 400  
building tree 50 of 400  
building tree 51 of 400  
building tree 52 of 400  
building tree 53 of 400  
building tree 54 of 400  
building tree 55 of 400  
building tree 56 of 400  
building tree 57 of 400  
building tree 58 of 400  
building tree 59 of 400  
building tree 60 of 400  
building tree 61 of 400  
building tree 62 of 400  
building tree 63 of 400  
building tree 64 of 400  
building tree 65 of 400  
building tree 66 of 400  
building tree 67 of 400  
building tree 68 of 400  
building tree 69 of 400

building tree 70 of 400  
building tree 71 of 400  
building tree 72 of 400  
building tree 73 of 400  
building tree 74 of 400  
building tree 75 of 400  
building tree 76 of 400  
building tree 77 of 400  
building tree 78 of 400  
building tree 79 of 400  
building tree 80 of 400  
building tree 81 of 400  
building tree 82 of 400  
building tree 83 of 400  
building tree 84 of 400  
building tree 85 of 400  
building tree 86 of 400  
building tree 87 of 400  
building tree 88 of 400  
building tree 89 of 400  
building tree 90 of 400  
building tree 91 of 400  
building tree 92 of 400  
building tree 93 of 400  
building tree 94 of 400  
building tree 95 of 400  
building tree 96 of 400  
building tree 97 of 400  
building tree 98 of 400  
building tree 99 of 400  
building tree 100 of 400  
building tree 101 of 400  
building tree 102 of 400  
building tree 103 of 400  
building tree 104 of 400  
building tree 105 of 400  
building tree 106 of 400  
building tree 107 of 400  
building tree 108 of 400  
building tree 109 of 400  
building tree 110 of 400  
building tree 111 of 400  
building tree 112 of 400  
building tree 113 of 400  
building tree 114 of 400  
building tree 115 of 400  
building tree 116 of 400  
building tree 117 of 400

building tree 118 of 400  
building tree 119 of 400  
building tree 120 of 400  
building tree 121 of 400  
building tree 122 of 400  
building tree 123 of 400  
building tree 124 of 400  
building tree 125 of 400  
building tree 126 of 400  
building tree 127 of 400  
building tree 128 of 400  
building tree 129 of 400  
building tree 130 of 400  
building tree 131 of 400  
building tree 132 of 400  
building tree 133 of 400  
building tree 134 of 400  
building tree 135 of 400  
building tree 136 of 400  
building tree 137 of 400  
building tree 138 of 400  
building tree 139 of 400  
building tree 140 of 400  
building tree 141 of 400  
building tree 142 of 400  
building tree 143 of 400  
building tree 144 of 400  
building tree 145 of 400  
building tree 146 of 400  
building tree 147 of 400  
building tree 148 of 400  
building tree 149 of 400  
building tree 150 of 400  
building tree 151 of 400  
building tree 152 of 400  
building tree 153 of 400  
building tree 154 of 400  
building tree 155 of 400  
building tree 156 of 400  
building tree 157 of 400  
building tree 158 of 400  
building tree 159 of 400  
building tree 160 of 400  
building tree 161 of 400  
building tree 162 of 400  
building tree 163 of 400  
building tree 164 of 400  
building tree 165 of 400

building tree 166 of 400  
building tree 167 of 400  
building tree 168 of 400  
building tree 169 of 400  
building tree 170 of 400  
building tree 171 of 400  
building tree 172 of 400  
building tree 173 of 400  
building tree 174 of 400  
building tree 175 of 400  
building tree 176 of 400  
building tree 177 of 400  
building tree 178 of 400  
building tree 179 of 400  
building tree 180 of 400  
building tree 181 of 400  
building tree 182 of 400  
building tree 183 of 400  
building tree 184 of 400  
building tree 185 of 400  
building tree 186 of 400  
building tree 187 of 400  
building tree 188 of 400  
building tree 189 of 400  
building tree 190 of 400  
building tree 191 of 400  
building tree 192 of 400  
building tree 193 of 400  
building tree 194 of 400  
building tree 195 of 400  
building tree 196 of 400  
building tree 197 of 400  
building tree 198 of 400  
building tree 199 of 400  
building tree 200 of 400  
building tree 201 of 400  
building tree 202 of 400  
building tree 203 of 400  
building tree 204 of 400  
building tree 205 of 400  
building tree 206 of 400  
building tree 207 of 400  
building tree 208 of 400  
building tree 209 of 400  
building tree 210 of 400  
building tree 211 of 400  
building tree 212 of 400  
building tree 213 of 400

building tree 214 of 400  
building tree 215 of 400  
building tree 216 of 400  
building tree 217 of 400  
building tree 218 of 400  
building tree 219 of 400  
building tree 220 of 400  
building tree 221 of 400  
building tree 222 of 400  
building tree 223 of 400  
building tree 224 of 400  
building tree 225 of 400  
building tree 226 of 400  
building tree 227 of 400  
building tree 228 of 400  
building tree 229 of 400  
building tree 230 of 400  
building tree 231 of 400  
building tree 232 of 400  
building tree 233 of 400  
building tree 234 of 400  
building tree 235 of 400  
building tree 236 of 400  
building tree 237 of 400  
building tree 238 of 400  
building tree 239 of 400  
building tree 240 of 400  
building tree 241 of 400  
building tree 242 of 400  
building tree 243 of 400  
building tree 244 of 400  
building tree 245 of 400  
building tree 246 of 400  
building tree 247 of 400  
building tree 248 of 400  
building tree 249 of 400  
building tree 250 of 400  
building tree 251 of 400  
building tree 252 of 400  
building tree 253 of 400  
building tree 254 of 400  
building tree 255 of 400  
building tree 256 of 400  
building tree 257 of 400  
building tree 258 of 400  
building tree 259 of 400  
building tree 260 of 400  
building tree 261 of 400

building tree 262 of 400  
building tree 263 of 400  
building tree 264 of 400  
building tree 265 of 400  
building tree 266 of 400  
building tree 267 of 400  
building tree 268 of 400  
building tree 269 of 400  
building tree 270 of 400  
building tree 271 of 400  
building tree 272 of 400  
building tree 273 of 400  
building tree 274 of 400  
building tree 275 of 400  
building tree 276 of 400  
building tree 277 of 400  
building tree 278 of 400  
building tree 279 of 400  
building tree 280 of 400  
building tree 281 of 400  
building tree 282 of 400  
building tree 283 of 400  
building tree 284 of 400  
building tree 285 of 400  
building tree 286 of 400  
building tree 287 of 400  
building tree 288 of 400  
building tree 289 of 400  
building tree 290 of 400  
building tree 291 of 400  
building tree 292 of 400  
building tree 293 of 400  
building tree 294 of 400  
building tree 295 of 400  
building tree 296 of 400  
building tree 297 of 400  
building tree 298 of 400  
building tree 299 of 400  
building tree 300 of 400  
building tree 301 of 400  
building tree 302 of 400  
building tree 303 of 400  
building tree 304 of 400  
building tree 305 of 400  
building tree 306 of 400  
building tree 307 of 400  
building tree 308 of 400  
building tree 309 of 400

building tree 310 of 400  
building tree 311 of 400  
building tree 312 of 400  
building tree 313 of 400  
building tree 314 of 400  
building tree 315 of 400  
building tree 316 of 400  
building tree 317 of 400  
building tree 318 of 400  
building tree 319 of 400  
building tree 320 of 400  
building tree 321 of 400  
building tree 322 of 400  
building tree 323 of 400  
building tree 324 of 400  
building tree 325 of 400  
building tree 326 of 400  
building tree 327 of 400  
building tree 328 of 400  
building tree 329 of 400  
building tree 330 of 400  
building tree 331 of 400  
building tree 332 of 400  
building tree 333 of 400  
building tree 334 of 400  
building tree 335 of 400  
building tree 336 of 400  
building tree 337 of 400  
building tree 338 of 400  
building tree 339 of 400  
building tree 340 of 400  
building tree 341 of 400  
building tree 342 of 400  
building tree 343 of 400  
building tree 344 of 400  
building tree 345 of 400  
building tree 346 of 400  
building tree 347 of 400  
building tree 348 of 400  
building tree 349 of 400  
building tree 350 of 400  
building tree 351 of 400  
building tree 352 of 400  
building tree 353 of 400  
building tree 354 of 400  
building tree 355 of 400  
building tree 356 of 400  
building tree 357 of 400

building tree 358 of 400  
building tree 359 of 400  
building tree 360 of 400  
building tree 361 of 400  
building tree 362 of 400  
building tree 363 of 400  
building tree 364 of 400  
building tree 365 of 400  
building tree 366 of 400  
building tree 367 of 400  
building tree 368 of 400  
building tree 369 of 400  
building tree 370 of 400  
building tree 371 of 400  
building tree 372 of 400  
building tree 373 of 400  
building tree 374 of 400  
building tree 375 of 400  
building tree 376 of 400  
building tree 377 of 400  
building tree 378 of 400  
building tree 379 of 400  
building tree 380 of 400  
building tree 381 of 400  
building tree 382 of 400  
building tree 383 of 400  
building tree 384 of 400  
building tree 385 of 400  
building tree 386 of 400  
building tree 387 of 400  
building tree 388 of 400  
building tree 389 of 400  
building tree 390 of 400  
building tree 391 of 400  
building tree 392 of 400  
building tree 393 of 400  
building tree 394 of 400  
building tree 395 of 400  
building tree 396 of 400  
building tree 397 of 400  
building tree 398 of 400  
building tree 399 of 400  
building tree 400 of 400

[Parallel(n\_jobs=1)]: Done 400 out of 400 | elapsed: 1.7s finished

[45]: RandomForestClassifier(max\_depth=5, max\_features=8, min\_samples\_leaf=7,  
n\_estimators=400, random\_state=123, verbose=2)



```
[46]: # Predictions for the training and testing sets
rf_test_pred = random_forest_model.predict(X_test)
rf_train_pred = random_forest_model.predict(X_train)

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=1)]: Done 400 out of 400 | elapsed: 0.0s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=1)]: Done 400 out of 400 | elapsed: 0.1s finished
```

```
[47]: # Accuracy scores for the training and testing sets
acc_rf_train= accuracy_score(y_train, rf_train_pred)
acc_rf_test = accuracy_score(y_test, rf_test_pred)

print("Train set Accuracy: %f " % (acc_rf_train))
print("Test set Accuracy: %f " % (acc_rf_test))

# AUC scores for the training and testing sets
roc_rf_train = roc_auc_score(y_train, rf_train_pred)
roc_rf_test = roc_auc_score(y_test, rf_test_pred)

print("Train set AUC: %f" % (roc_rf_train))
print("Test set AUC: %f" % (roc_rf_test))

# F1 scores for the training and testing sets
f1_rf_train = f1_score(y_train, rf_train_pred, average='macro')
f1_rf_test = f1_score(y_test, rf_test_pred, average='macro')

print("Train set F1: %f" % (f1_rf_train))
print("Test set F1: %f" % (f1_rf_test))
```

```
Train set Accuracy: 0.774656
Test set Accuracy: 0.720824
Train set AUC: 0.756960
Test set AUC: 0.697867
Train set F1: 0.761811
Test set F1: 0.701570
```

```
[48]: # Create a DataFrame to store the evaluation results
metrics_rf_df = pd.DataFrame({
    'Train': [acc_rf_train, roc_rf_train, f1_rf_train],
    'Test': [acc_rf_test, roc_rf_test, f1_rf_test]
}, index=['Accuracy', 'AUC', 'F1 Score'])

# Print the DataFrame
print(metrics_rf_df)
```

	Train	Test
Accuracy	0.774656	0.720824
AUC	0.756960	0.697867
F1 Score	0.761811	0.701570

```
[49]: # Classification report for the best XGBoost model
print(classification_report(y_test, rf_test_pred))
```

	precision	recall	f1-score	support
0	0.73	0.84	0.78	255
1	0.71	0.56	0.63	182
accuracy			0.72	437
macro avg	0.72	0.70	0.70	437
weighted avg	0.72	0.72	0.71	437

When we run the model without hyperparameter tuning, it does worse. Hyperparameter tuning is essential to ensure our model is making the best choices to avoid overfitting.

We can display the most important features:

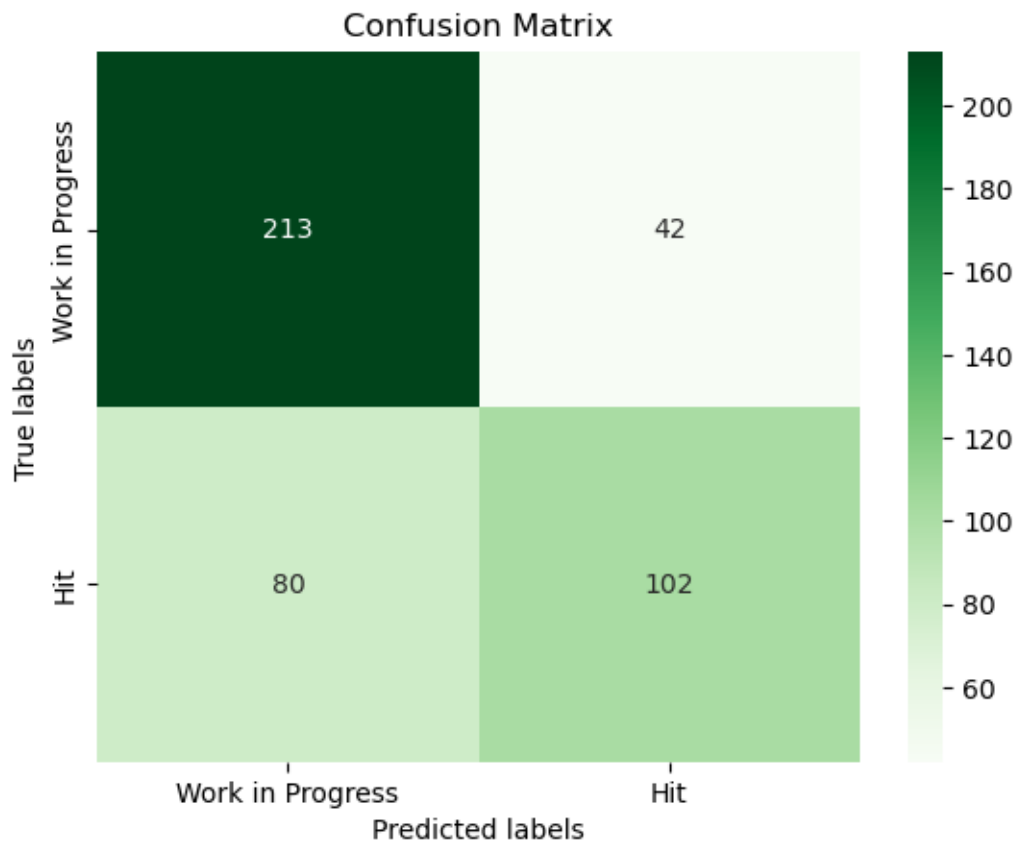
```
[50]: pd.Series(random_forest_model.feature_importances_, index = X.columns).
      ↪sort_values(ascending = False)
```

```
[50]: artist_popularity      0.417694
year                        0.306164
popularity_difference      0.044037
acousticness               0.042393
loudness                   0.033851
duration_ms               0.033671
Melodic Energy             0.028978
Emotional Tempo            0.028901
speechiness                0.027978
liveness                   0.026114
instrumentalness            0.010218
dtype: float64
```

```
[51]: # Initialize the confusion matrix
cm = confusion_matrix(y_test, rf_test_pred)
ax= plt.subplot()
# Create a heatmap of the confusion matrix
sns.heatmap(cm, annot=True, ax = ax, fmt='g', cmap='Greens'); #annot=True to
    ↪annotate cells

# Setting labels, title and ticks
ax.set_xlabel('Predicted labels');ax.set_ylabel('True labels');
ax.set_title('Confusion Matrix');
```

```
ax.xaxis.set_ticklabels(['Work in Progress', 'Hit']); ax.yaxis.  
    ↪set_ticklabels(['Work in Progress', 'Hit']);
```



## 0.5 XGBoost model

Let's attempt an XGBoost Regressor to see if we can get a better performance.

```
[68]: # Instantiate XGBoost model
model = XGBClassifier(n_estimators=200)
# Fit the model to the training data
model.fit(X_train, y_train)
scorer = 'accuracy'
# Parameters for GridSearchCV for XGBoost
params = {
    'subsample': [0.6, 0.8, 1.0], # the fraction of samples to be used for
    ↪fitting the individual base learners
    'max_depth': [3, 5, 7], # maximum number of levels in each decision tree
    'learning_rate': [0.001, 0.01, 0.1] # step size shrinkage used to
    ↪prevent overfitting
}
```

```

# Instantiate the grid search model with the XGBoost model and the parameter
↳grid
xgbGrid = GridSearchCV(model, params, cv = 5, scoring='roc_auc', verbose=1)
# Fit the grid search to the data
xgbGrid.fit(X_train, y_train)

```

Fitting 5 folds for each of 27 candidates, totalling 135 fits

```

[68]: GridSearchCV(cv=5,
                  estimator=XGBClassifier(base_score=None, booster=None,
                                           callbacks=None, colsample_bylevel=None,
                                           colsample_bynode=None,
                                           colsample_bytree=None, device=None,
                                           early_stopping_rounds=None,
                                           enable_categorical=False, eval_metric=None,
                                           feature_types=None, gamma=None,
                                           grow_policy=None, importance_type=None,
                                           interaction_constraints=None,
                                           learning_rate=None, ...
                                           max_cat_threshold=None,
                                           max_cat_to_onehot=None,
                                           max_delta_step=None, max_depth=None,
                                           max_leaves=None, min_child_weight=None,
                                           missing=nan, monotone_constraints=None,
                                           multi_strategy=None, n_estimators=200,
                                           n_jobs=None, num_parallel_tree=None,
                                           random_state=None, ...),
                  param_grid={'learning_rate': [0.001, 0.01, 0.1],
                              'max_depth': [3, 5, 7], 'subsample': [0.6, 0.8, 1.0]},
                  scoring='roc_auc', verbose=1)

```

```

[69]: # fitting grid search to XGB model
XGB_model = XGBClassifier(n_estimators = 200, random_state=123,
                          subsample=xgbGrid.best_params_['subsample'], # from
↳our grid search
                          max_depth=xgbGrid.best_params_['max_depth'], # from
↳our grid search
                          learning_rate=xgbGrid.
↳best_params_['learning_rate'], # from our grid search
                          verbose=2) # verbose = 2 means that we want to see
↳the output
XGB_model.fit(X_train, y_train)
# pred_train = XGB_model.predict(X_train)
# pred_test = XGB_model.predict(X_test)

```

/Users/nataliaclark/anaconda3/lib/python3.11/site-packages/xgboost/core.py:160:

UserWarning:

[21:38:47] WARNING: /Users/runner/work/xgboost/xgboost/src/learner.cc:742:  
Parameters: { "verbose" } are not used.

```
[69]: XGBClassifier(base_score=None, booster=None, callbacks=None,  
                  colsample_bylevel=None, colsample_bynode=None,  
                  colsample_bytree=None, device=None, early_stopping_rounds=None,  
                  enable_categorical=False, eval_metric=None, feature_types=None,  
                  gamma=None, grow_policy=None, importance_type=None,  
                  interaction_constraints=None, learning_rate=0.01, max_bin=None,  
                  max_cat_threshold=None, max_cat_to_onehot=None,  
                  max_delta_step=None, max_depth=3, max_leaves=None,  
                  min_child_weight=None, missing=nan, monotone_constraints=None,  
                  multi_strategy=None, n_estimators=200, n_jobs=None,  
                  num_parallel_tree=None, random_state=123, ...)
```

```
[70]: # Prediction using XGBoost model with the best parameters found by GridSearchCV  
pred_train = XGB_model.predict_proba(X_train)[:, 1]  
pred_test = XGB_model.predict_proba(X_test)[:, 1]  
  
# Convert predicted probabilities to binary predictions (0 or 1) using a  
↪ threshold of 0.5  
pred_train_binary = (pred_train >= 0.5).astype(int)  
pred_test_binary = (pred_test >= 0.5).astype(int)
```

```
[71]: # Accuracy scores for the training and testing sets  
acc_boost_train = accuracy_score(y_train, pred_train_binary)  
acc_boost_test = accuracy_score(y_test, pred_test_binary)  
  
print("Attempt 1 Train Accuracy: %f" % (acc_boost_train))  
print("Attempt 1 Test Accuracy: %f" % (acc_boost_test))  
  
# AUC scores for the training and testing sets  
roc_boost_train = roc_auc_score(y_train, pred_train)  
roc_boost_test = roc_auc_score(y_test, pred_test)  
  
print("Train set AUC: %f" % (roc_boost_train))  
print("Test set AUC: %f" % (roc_boost_test))  
  
f1_boost_train = f1_score(y_train, pred_train_binary)  
f1_boost_test = f1_score(y_test, pred_test_binary)  
print("F1 Score Train: ", f1_boost_train)  
print("F1 Score Test: ", f1_boost_test)
```

Attempt 1 Train Accuracy: 0.749427

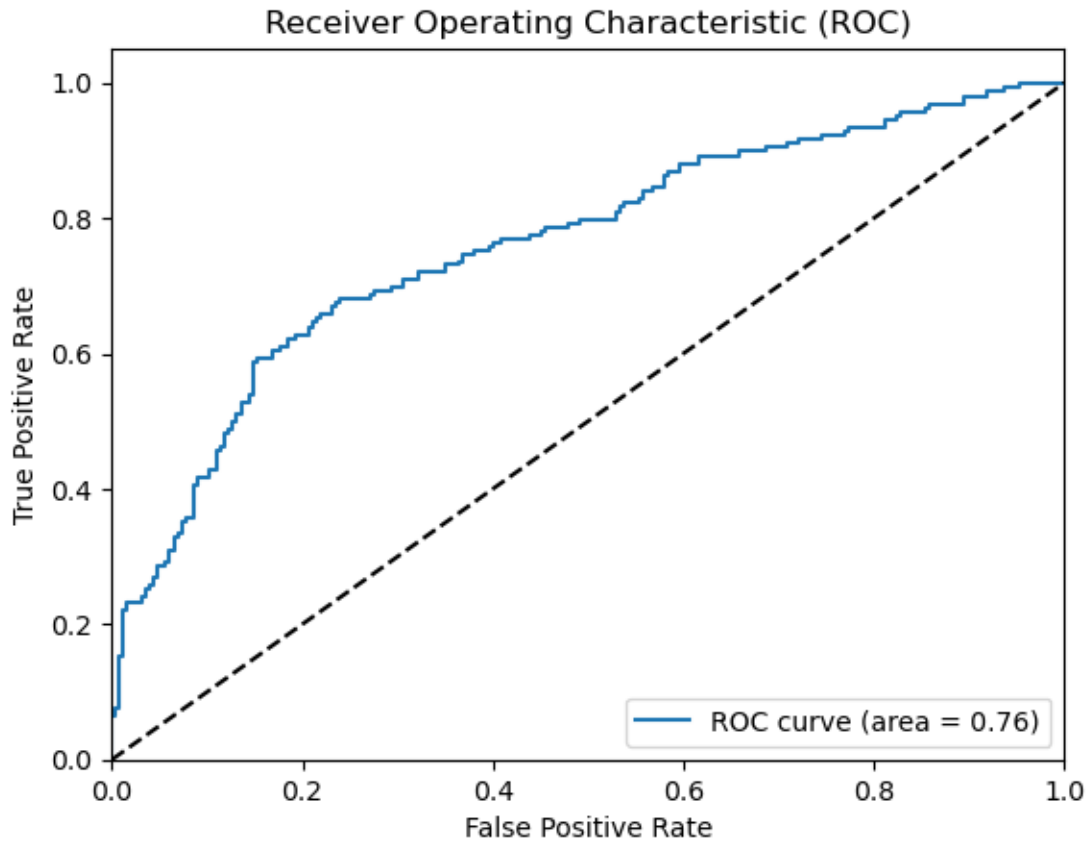
```
Attempt 1 Test Accuracy: 0.720824
Train set AUC: 0.824440
Test set AUC: 0.760267
F1 Score Train: 0.6630686198920586
F1 Score Test: 0.6139240506329114
```

```
[72]: # Create a DataFrame to store the evaluation results
metrics_boost_df = pd.DataFrame({
    'Train': [acc_boost_train, roc_boost_train, f1_boost_train],
    'Test': [acc_boost_test, roc_boost_test, f1_boost_test]
}, index=['Accuracy', 'AUC', 'F1'])

# Print the DataFrame
print(metrics_boost_df)
```

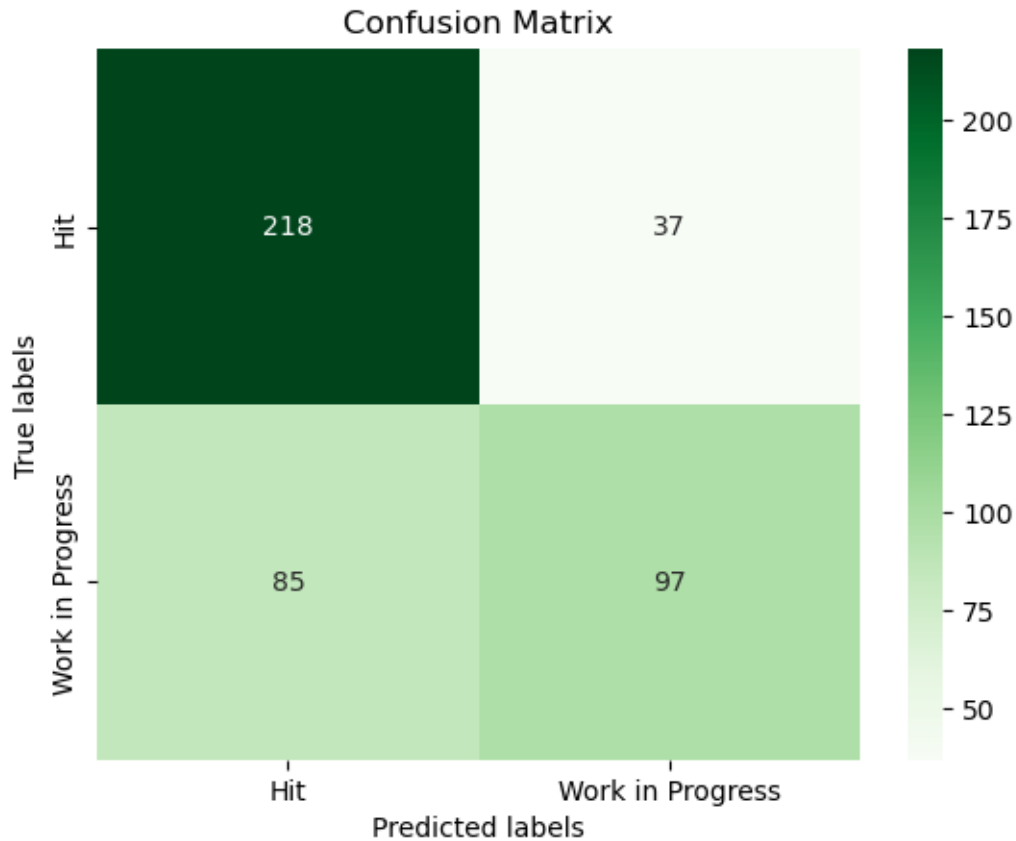
	Train	Test
Accuracy	0.749427	0.720824
AUC	0.824440	0.760267
F1	0.663069	0.613924

```
[57]: # Plot the ROC curve
# Initialize the plot
fpr, tpr, thresholds = roc_curve(y_test, pred_test)
# Calculate the AUC
plt.plot(fpr, tpr, label='ROC curve (area = %0.2f)' % roc_boost_test)
# Plot the diagonal line from (0, 1)
plt.plot([0, 1], [0, 1], 'k--')
# Set the limits of the plot for x and y axes
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
# Set the labels and title
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
# Show the legend
plt.legend(loc="lower right")
# Show the plot
plt.show()
```



```
[58]: # Initialize the confusion matrix
cm = confusion_matrix(y_test, pred_test_binary)
ax= plt.subplot()
# Create a heatmap of the confusion matrix
sns.heatmap(cm, annot=True, ax = ax, fmt='g', cmap='Greens'); #annot=True to
    ↪ annotate cells

# Setting labels, title and ticks
ax.set_xlabel('Predicted labels');ax.set_ylabel('True labels');
ax.set_title('Confusion Matrix');
ax.xaxis.set_ticklabels(['Hit', 'Work in Progress']); ax.yaxis.
    ↪set_ticklabels(['Hit', 'Work in Progress']);
```



```
[59]: # Finding feature importances from the best XGBoost model
feature_importances = XGB_model.feature_importances_

# Create a Series with feature importances as values and feature names as index
feature_importance_series = pd.Series(feature_importances, index=X.columns)

# Sort the Series by importance values in descending order
feature_importance_series_sorted = feature_importance_series.
    ↪sort_values(ascending=False)

# Show the feature importances
print(feature_importance_series_sorted)
```

```
artist_popularity    0.239313
year                 0.231101
popularity_difference 0.084452
loudness             0.071111
duration_ms          0.066785
Melodic Energy       0.057306
acousticness         0.054091
```



```
speechiness          0.050549
Emotional Tempo      0.049564
liveness             0.048824
instrumentalness      0.046905
dtype: float32
```

```
[60]: # Classification report for the best XGBoost model
print(classification_report(y_test, pred_test_binary))
```

	precision	recall	f1-score	support
0	0.72	0.85	0.78	255
1	0.72	0.53	0.61	182
accuracy			0.72	437
macro avg	0.72	0.69	0.70	437
weighted avg	0.72	0.72	0.71	437

### 0.5.1 Appendix: Considerations to follow thinking process

- Genre variable: it was referring to the artist, not the track itself. We considered a segmentation of the variable into groups, but when we tried with these variables, they did not add information to our model.
- Choosing a model: our data had outliers and not too many features, made sense to use a tree based model. We tried a simple Decision Tree, Bagging, Random Forest, and Boosting. Boosting seemed to consistently show us the best results, so we decided to use it. We have shown the code here for our 2 most successful models, seeing as we wanted to reduce the complexity of the notebook.
- Feature Engineering: How can we find more information in combinations of variables? We decided that danceability and energy could be multiplied, valence and tempo could be multiplied, and we could come up with a feature for the difference of the percentage of songs made by “popular artists” to the track’s artists percentage of songs made. We set this threshold at the 75th percentile.
- What do we do with outliers?: Since we were deciding to use a tree-based model, we knew that outliers are implicitly handled well with these models. We did not worry about clipping them.
- Train/test split: 2300 records is not a lot - we want to have enough training data, but also reserve enough for testing.
- Overfitting: number of trees doesn’t impact with most of these models, so we focused on the hyperparameters that would cause our model to overfit for each model and tried to tune on those.
- Computational complexity: considered, which is why we limited grid search parameters using logic and best judgment
- Feature importance: Seems to be artist popularity for all models tested. How can we turn this into a business approach to our problem?
- How do we choose an evaluation metric that best encapsulates our goal? How do we defend the accuracy numbers without undermining our model?